

# CS 186 Section 4: Buffers, Files, and Indexes


Vikram Sreekanti

# Buffer Replacement

- Buffer pool is your memory manager.
  - Data must be in memory to operate on it; buffer manager hides the fact that not all data is in memory.
    - This should sound familiar to everyone (especially anyone who's taken 162).
- We've said it a million times: your data doesn't fit in memory.
- What do you do when memory is full, and you want to load something else in?
  - You kick something out!
  - There are many, many different buffer replacement algorithms.
  - Common algorithms:
    - LRU
    - MRU
    - Clock
    - etc.


# Buffer Replacement: Least Recently Used

Main Idea: When you need to evict something, you evict the page that was least-recently used.




# Buffer Replacement: Most Recently Used

Main Idea: When you need to evict something, you evict the page that was most-recently used.



# Buffer Replacement: Clock

Main Idea: Close approximation of LRU, but each page gets a second chance bit.



# Buffer Replacement Exercises



You have four buffer pages. What is the end state of the buffer pool for {LRU, MRU, Clock} after the following access pattern? What is your hit rate?

**A B C D A F A D G D G E D F**

# Buffer Replacement

## Exercise: LRU

**A B C D A F A D G D G E D F**




Number of Hits:

0

# Buffer Replacement

## Exercise: MRU

**A B C D A F A D G D G E D F**



Number of Hits:



0



# Buffer Replacement

## Exercise: Clock

**A B C D A F A D G D G E D F**



Number of Hits:

0

Clock Hand:

A

# Heap (Unsorted) File

IMPORTANT: Do not confuse this with the heap data structure — there no relation.

	Heap File (IOs)
Sequential Scan	B
Equality Search	.5B
Range Search	B
Insertion	2
Deletion	.5B + 1

# Sorted File

	Heap File (IOs)	Sorted File (IOs)
Sequential Scan	$B$	$B$
Equality Search	$.5B$	$\log_2(B)$
Range Search	$B$	$\log_2(B) + \text{num\_matches}$
Insertion	$2$	$\log_2(B) + 2 * .5B$
Deletion	$.5B + 1$	(same as Insertion)

# Indexes (Indices?)

Disk-based data structure for fast value-based lookups of data in tables.

Indices come in many shapes and sizes: tree-based, hash; clustered vs. unclustered; alternatives 1, 2, and 3

# Three Storage Alternatives for Indexing

1. The pages of the index store the actual tuples.
2. The pages of the index store  $\langle \text{key}, \text{rid} \rangle$  pairs, where rid is one record with key.
3. The pages of the index store  $\langle \text{key}, \text{list}[\text{rid}] \rangle$  pairs, where list[rid] is all rids with key.
  - A. Can anyone see any issues with this?

# Clustered vs. Unclustered Indexes

- Clustered: Entries in the index are stored (approximately) in sorted order over key.
  - Alternative 1 indices are always clustered.
- Unclustered: ... not clustered.

Note: Clustering affects the way that the actual file is stored on disk.

# Index and Storage Exercises

What are important factors in determining whether or not you should add an index to a table?

Can't use hash indexing with range queries. Should know which field to cluster on (calculate I/Os based on typical queries that you'll need to run). Decide if you even want to cluster (high maintenance cost).

# Index and Storage Exercises

Consider the table:

`Enrolled(sid, course, grade)`

and the query:

```
SELECT * FROM Enrolled where sid > 4500;
```

Assume SIDs are unique and are in the range [1, 6000]



# Index and Join Exercises

How many IOs would this query take if the table was stored in a heap file?

$$B = 500$$

# Index and Join Exercises

How many IOs would this query take if the table was stored in a sorted file sorted on grades?

$$B = 500$$

# Index and Join Exercises

How many IOs would this query take if the table was stored in a sorted file sorted on SID?

$$\begin{aligned} B &= \log_2(500) + .25 (500) \\ &= 9 + 125 \\ &= 134 \text{ IOs} \end{aligned}$$

# Index and Join Exercises

**True or False?** Given the table `Students(sid, gpa, age)`, a hash index on GPA will significantly increase the performance of the following query:

```
SELECT * FROM Students WHERE age > 20;
```

False

# Index and Join Exercises

**True or False?** Given the table `Students(sid char(20), gpa float, age integer)`, a clustered tree based index on `gpa` will increase the performance of the following query:

```
SELECT * FROM Students where age > 20 AND  
gpa > 3.5;
```

True