

CS 186 Section 3: SQL, Joins, and Files

Vikram Sreekanti

Multi-Table SQL Queries

Multi-Table SQL

- **Main idea: These queries are very similar to single table queries.**
- Main differences:
 - There will be multiple tables listed in the `FROM` clause.
 - You (usually) specify a “join predicate”.
 - What is a join predicate?
 - A predicate on a value common to both tables.
 - e.g. `Students.sid = Enrollment.sid`
 - Where does it go?
 - The `WHERE` clause

Multi-Table SQL

Songs (song_id, song_name, album_num, weeks_in_top_40)

Artists(artist_id, artist_name, first_year_active)

Albums (album_id, album_name, artist_num, year_released,
genre)

Example Query

Songs (song_id, song_name, album_num, weeks_in_top_40)

Artists(artist_id, artist_name, first_year_active)

Albums (album_id, album_name, artist_num, year_released, genre)

```
SELECT song_name, album_name
FROM Songs S, Artists A
WHERE S.album_num = A.album_num;
```

SQL Exercises

```
Songs (song_id, song_name, album_num, weeks_in_top_40)
Artists(artist_id, artist_name, first_year_active)
Albums (album_id, album_name, artist_num, year_released,
genre)
```

1. The name of all songs with the genre “country” which have spent more than 2 weeks in the top 40.

```
SELECT Songs.song_name
FROM Albums, Songs
WHERE Songs.album_num = Albums.album_id
      AND Albums.genre = 'country'
      AND Songs.weeks_in_top_40 > 2;
```

SQL Exercises

```
Songs (song_id, song_name, album_num, weeks_in_top_40)
Artists(artist_id, artist_name, first_year_active)
Albums (album_id, album_name, artist_num, year_released,
genre)
```

2. For each song, its name, the name of its album, and the name of its artist.

```
SELECT Songs.song_name, Albums.album_name, Artists.artist_name
FROM Artists, Albums, Songs
WHERE Artists.artist_id = Albums.artist_num
      AND Songs.album_num = Albums.album_id;
```

SQL Exercises

Songs (song_id, song_name, album_num, weeks_in_top_40)
Artists(artist_id, artist_name, first_year_active)
Albums (album_id, album_name, artist_num, year_released,
genre)

3. The number of albums released by each artist.

```
SELECT count(*)  
FROM Artists, Albums  
WHERE Artists.artist_id = Albums.artist_num  
GROUP BY Artists.artist_id;
```


Join Algorithms

Cost Notation

- $[R]$ = the number of pages required to store the relation
- $|R|$ = the number of records in the relation
- PR = the number of records that fit on a single page of R

Simple Nested-Loops Join (SNLJ)

```
for record r in R:  
  for record s in S:  
    if theta(r, s):  
      add join(r, s) to result
```

COST: $IRI * [S] + [R]$

Page Nested-Loops Join (PNLJ)

```
for page p_r in R:  
  for page p_s in S:  
    for record r in p_r:  
      for record s in p_s:  
        if theta(r, s):  
          add join(r, s) to result
```

COST: $[R] * [S] + [R]$

Chunk Nested-Loops Join (CNLJ)

- This is very similar to PNLJ, but for one flaw win PNLJ.
 - PNLJ doesn't take advantage of all of memory.
- Instead of only loading *one* page of the outer relation into memory at a time, load B-2!

$$\text{COST: } ([R]/(B-2)) * [S] + [R]$$

Index Nested-Loops Join (INLJ)

```
for record r in R:  
  for record s in S where r == s:  
    add join(r, s) to result
```

This code is very similar to SNLJ, but the difference is that the second line does an index lookup.

Sort-Merge Join (SMJ)

Sort R on join attribute

Sort S on join attribute

Scan sorted-R and sorted-S in tandem to find matches

COST: $\text{cost}(\text{Sort}(R)) + \text{cost}(\text{Sort}(S)) + [R] + [S]$

Sort-Merge Join (SMJ)

Refinement:

Do the join during the final merging pass of sort.

Advantages of SMJ:

- one or both inputs is already sorted on join attribute
- output needs to be sorted on join attribute

COST: $\text{cost}(\text{Sort}(R)) + \text{cost}(\text{Sort}(S)) + [R] + [S]$

If you can sort in 2 passes: $3[R] + 3[S]$

Hash Join

```
externally hash R on join attribute  
externally hash S on join attribute  
for partition p_r in r:  
    build in memory hash table on a partition of R  
    stream s and probe in memory hash table
```

COST: $3([R] + [S]) + [\text{output}]$ (for 2 passes)

Join Exercises

Give scenarios in which each of the following is optimal:

- CNLJ
- SMJ
- Hash join

Join Exercises

Give scenarios in which each of the following is optimal:

- CNLJ
 - Not using an equality predicate!
 - Join is just a cartesian product; very few distinct keys
- SMJ
 - Skewed input data (excepting worst cases)
 - Small memory size
 - Want sorted output/have sorted input
- Hash join
 - One partition is large, the other is small
 - Hybrid hashing

Join Exercises

We have 12 pages of memory, and we want to join two tables R and S where [R] is 100 pages and [S] is 50 pages.

1. How many disk *reads* are done to perform CNLJ?
2. What about hash join?

Join Exercises

We have 12 pages of memory, and we want to join two tables R and S where [R] is 100 pages and [S] is 50 pages.

1. How many disk *reads* are done to perform CNLJ?

$$[S] + [S]/(B-2) * [R]$$

$$50 + (50/10) * (100) = 550$$

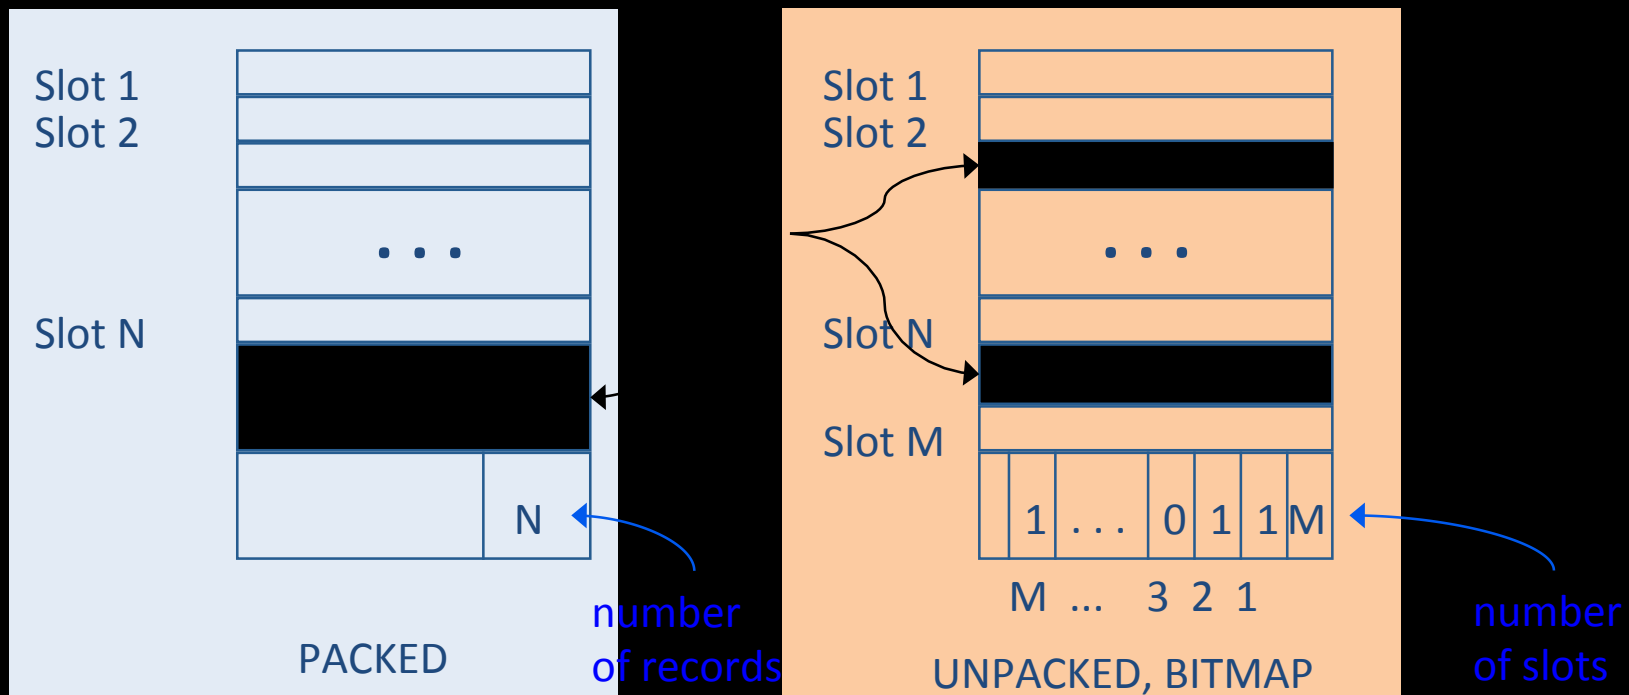
2. What about hash join?

$$([S] + [R]) * 2$$

$$150 * 2 = 300$$


Files and Pages

Pages: Fixed-Length Records



Record id = <page id, slot #>

Pages: Variable Length Records (Slotted Pages)



Pages & Files Exercises

What are the advantages and disadvantages of using slotted pages or bitmaps over just tightly packing records together?

Pages & Files Exercises

What are the advantages and disadvantages of using slotted pages or bitmaps over just tightly packing records together?

- allow movement of records without changing record ID
- slotted pages support variable-length records
- page directory incurs overhead - none necessary when using packed formats

Pages & Fields Exercises

You have a slotted page with 80 bytes of free space, and it costs 4 bytes to store a directory entry.

What is the size of largest record you can insert?

How many 1-byte sized records can you insert?

Pages & Fields Exercises

You have a slotted page with 80 bytes of free space, and it costs 4 bytes to store a directory entry.

What is the size of largest record you can insert?

A slot entry needs 4 bytes. $(80 - 4) = 76$ bytes.

How many 1-byte sized records can you insert?

Each record takes 1 byte for the record and 4 for the slot entry. That means we need 5 bytes per record. $80 \text{ bytes} / 5 \text{ bytes} = 16$ records.