

CS 186 Section 6: Advanced SQL, Data Modeling

Vikram Sreekanti

Data Modeling

ER Model

- Entity: A real-world object described by a set of attributes.
 - Entity Set (box): A collection of similar entities. (i.e., all Students.)
 - All entities in a set have the same attributes.
 - Entity sets have a primary key (underlined).
- Relationship: Association between two or more entities.
 - Relationship Set (diamond): collection of similar relationships.
 - Describe the interactions between entities.
 - Relationships can also have attributes.

Constraints





- An arrow denotes a key constraint. This means the key is there at most once.



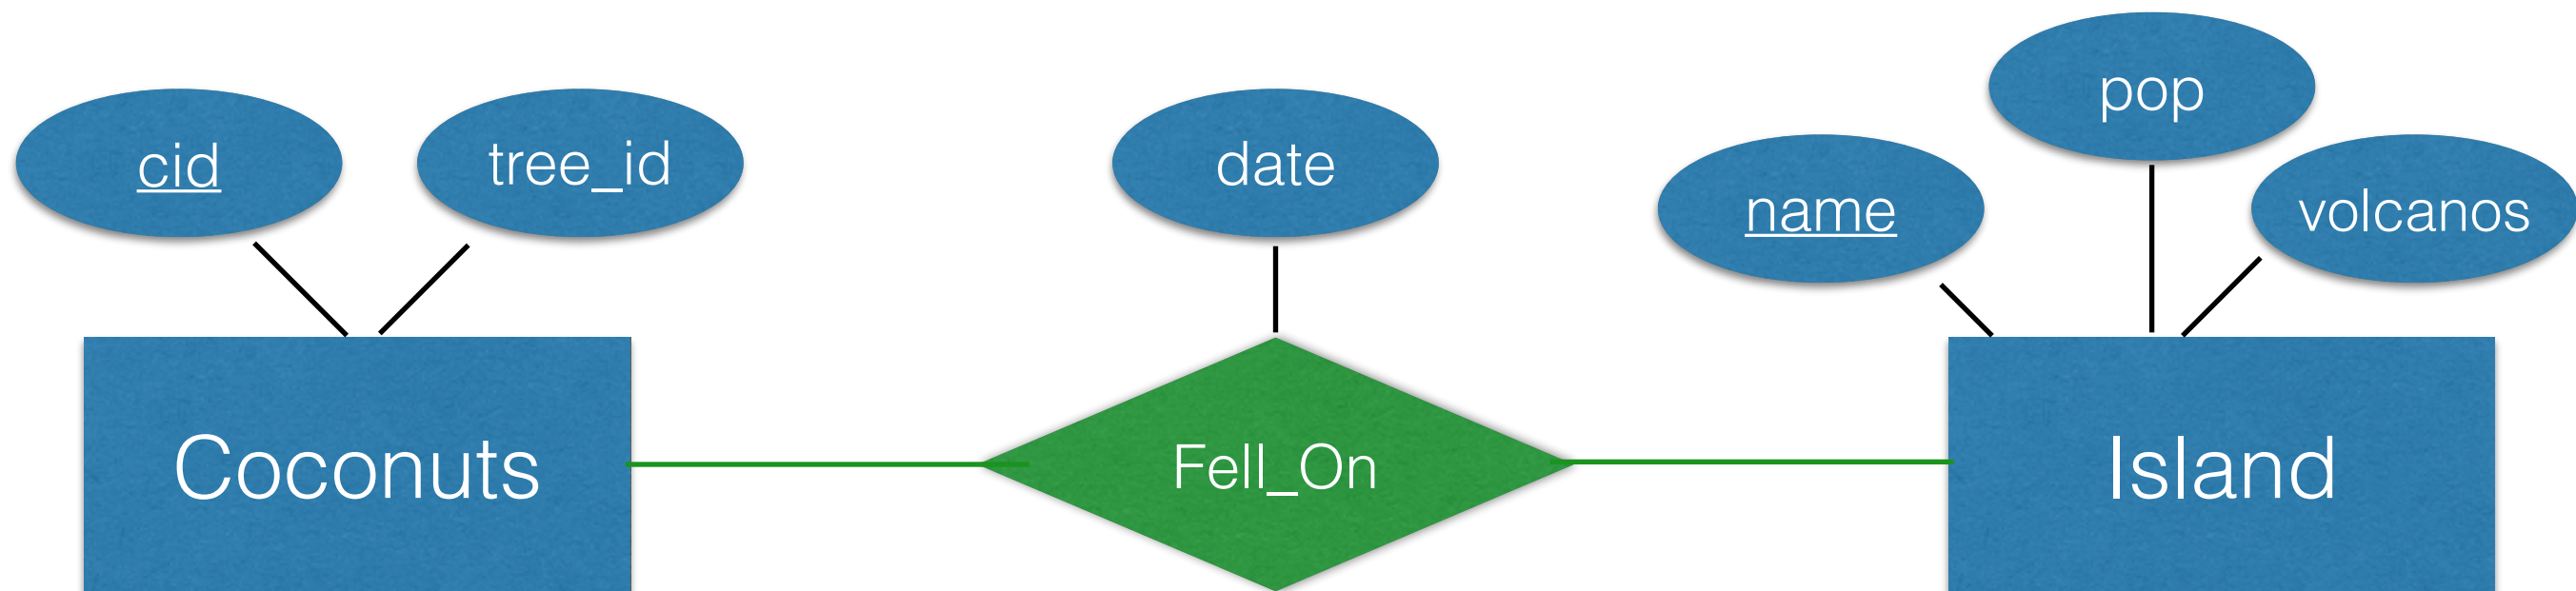
- A bold line denotes a participation constraint. This means “at least one”.



Constraints Explained Better

	Partial Participation	Total Participation
Non-Key	0 or more 	1 or more 
Key	0 or 1 	Exactly 1 

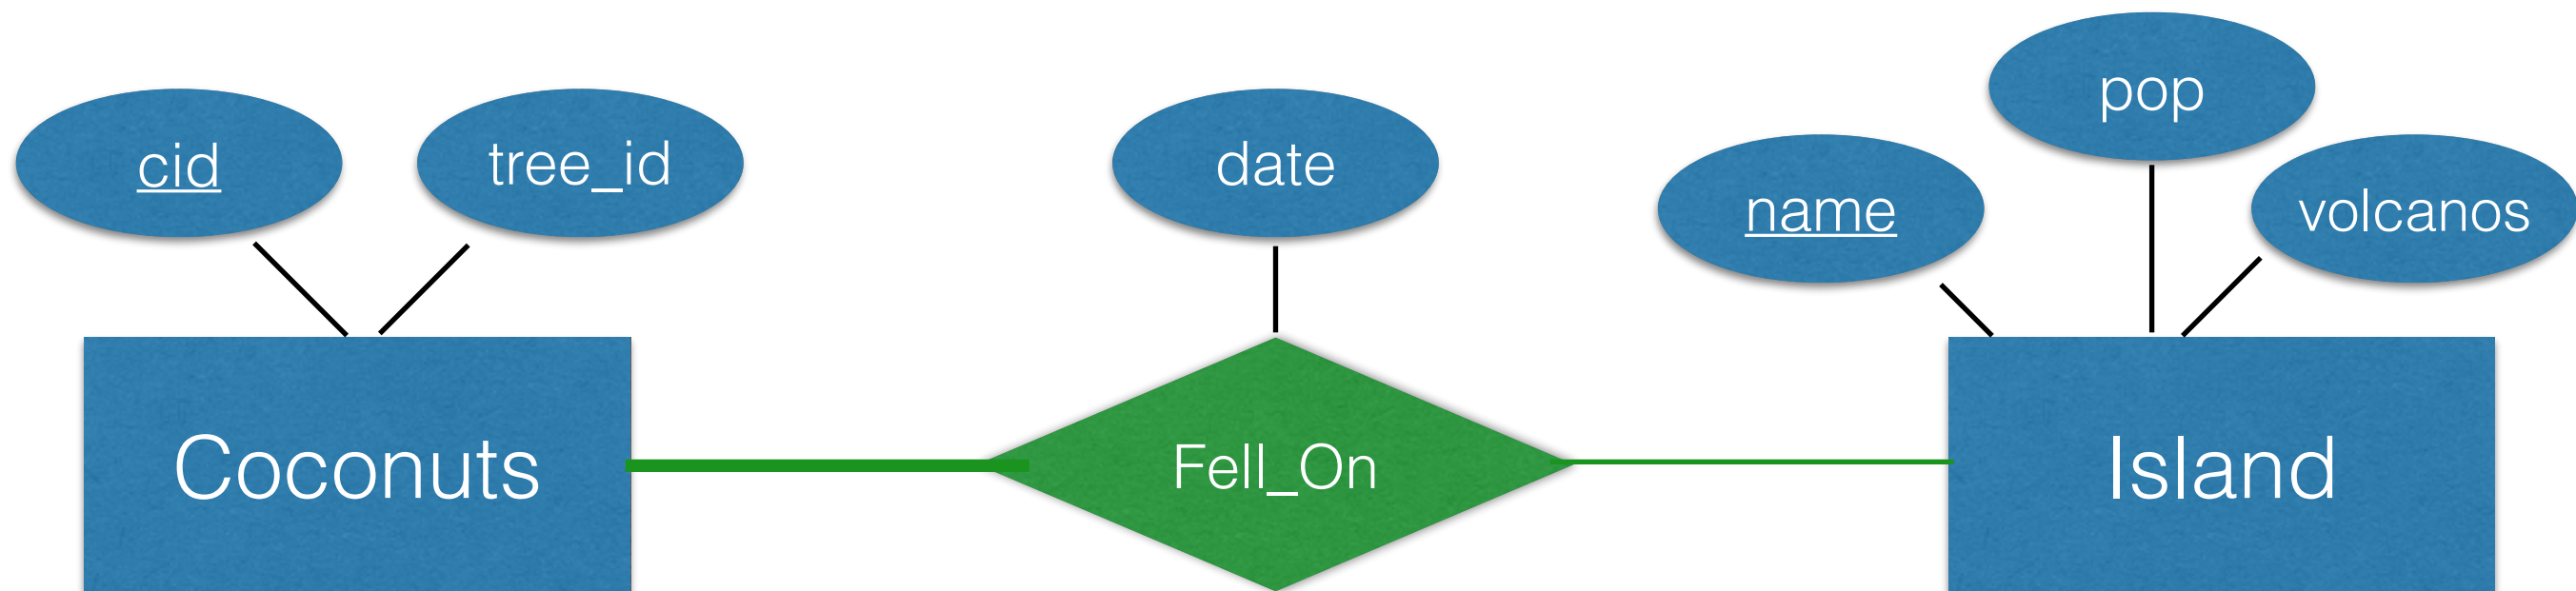
ER Diagrams: Quick Example



“Coconut participates in the Fell_On relationship 0 or more times.”
???

“A coconut can fall on an island zero or more times.”

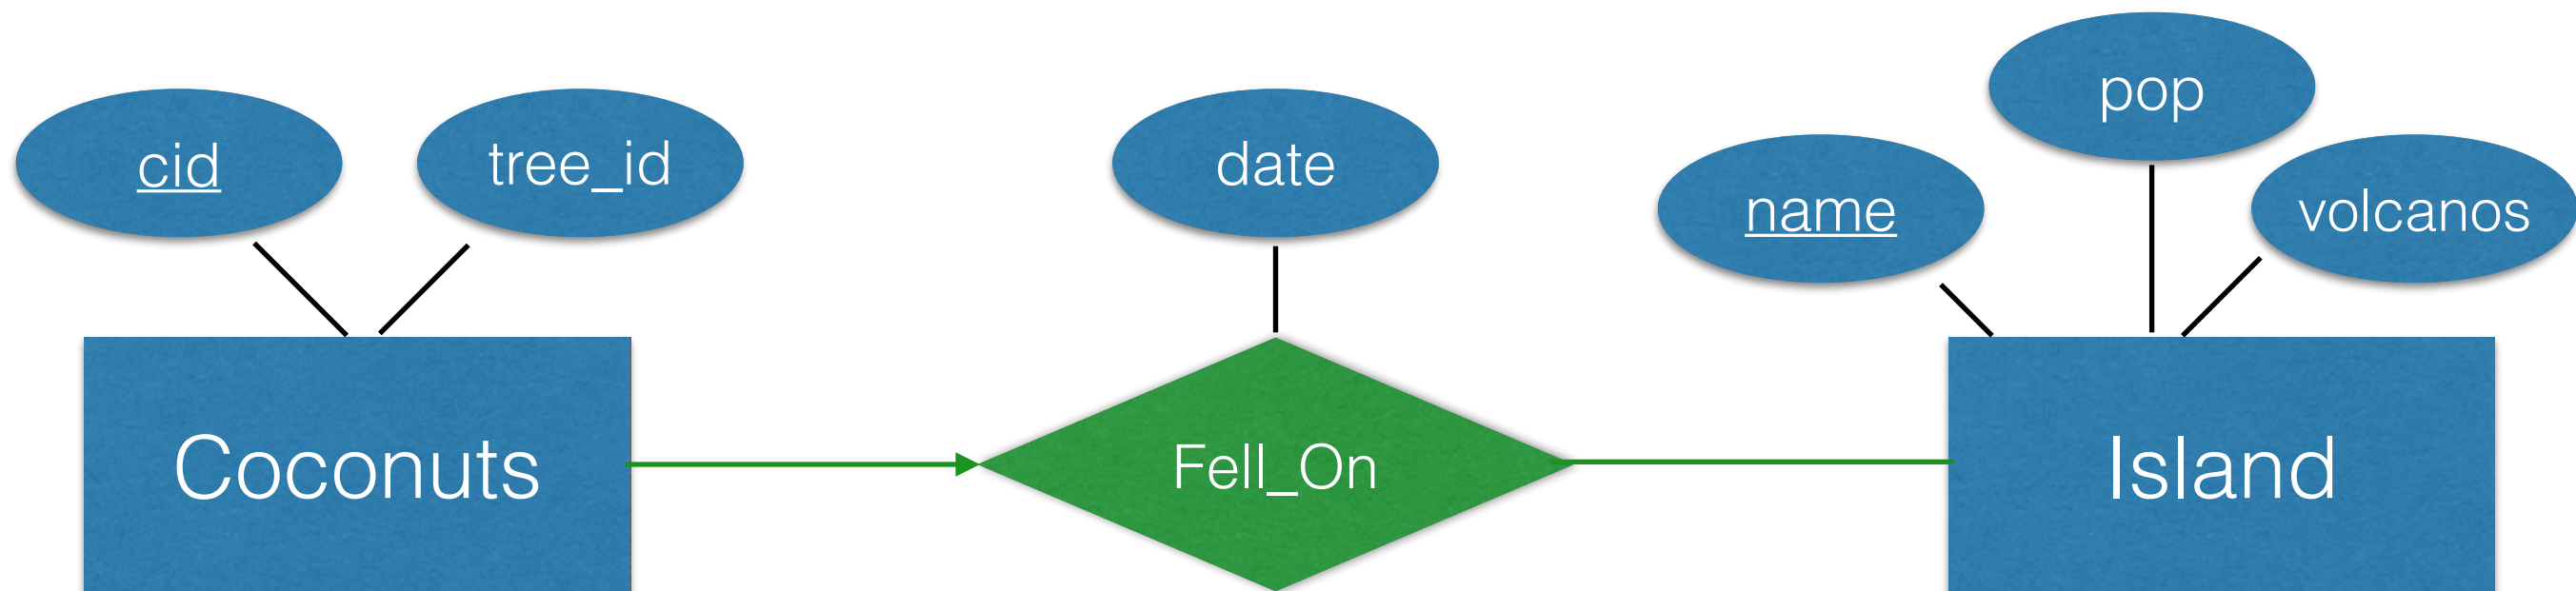
ER Diagrams: Quick Example



“Coconut participates in the Fell_On relationship 1 or more times.”
???

“A coconut must fall on an island at least once.”

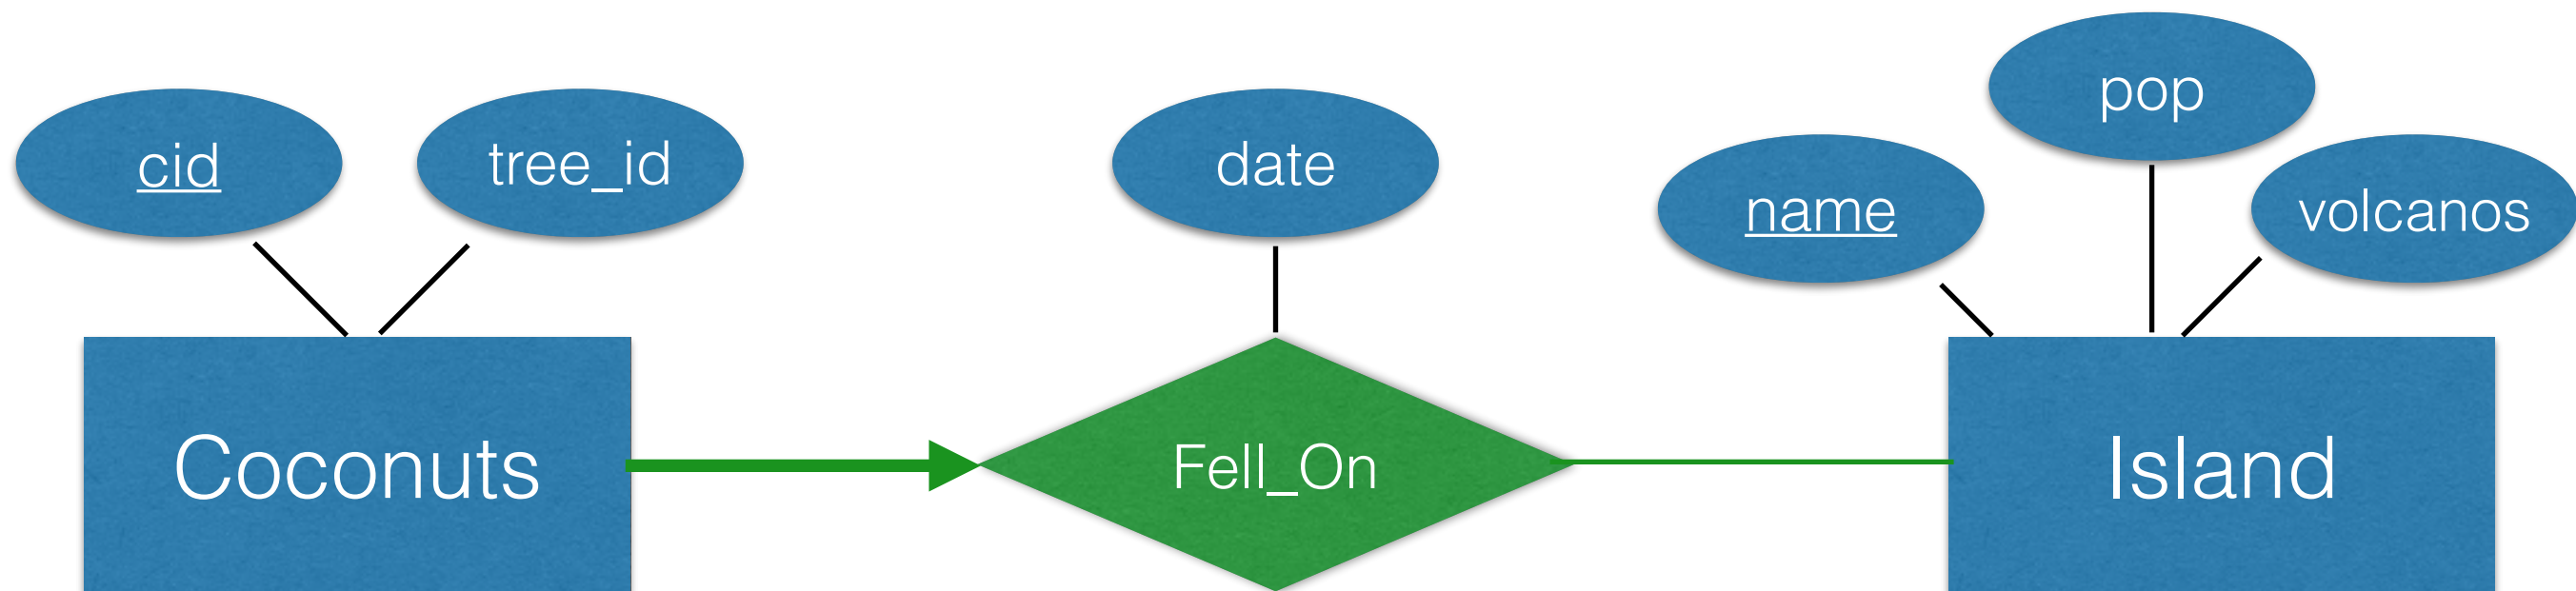
ER Diagrams: Quick Example



“Coconut participates in the Fell_On relationship 0 or 1 times.”
???

“A coconut can fall on an island at least 0 or 1 times.”

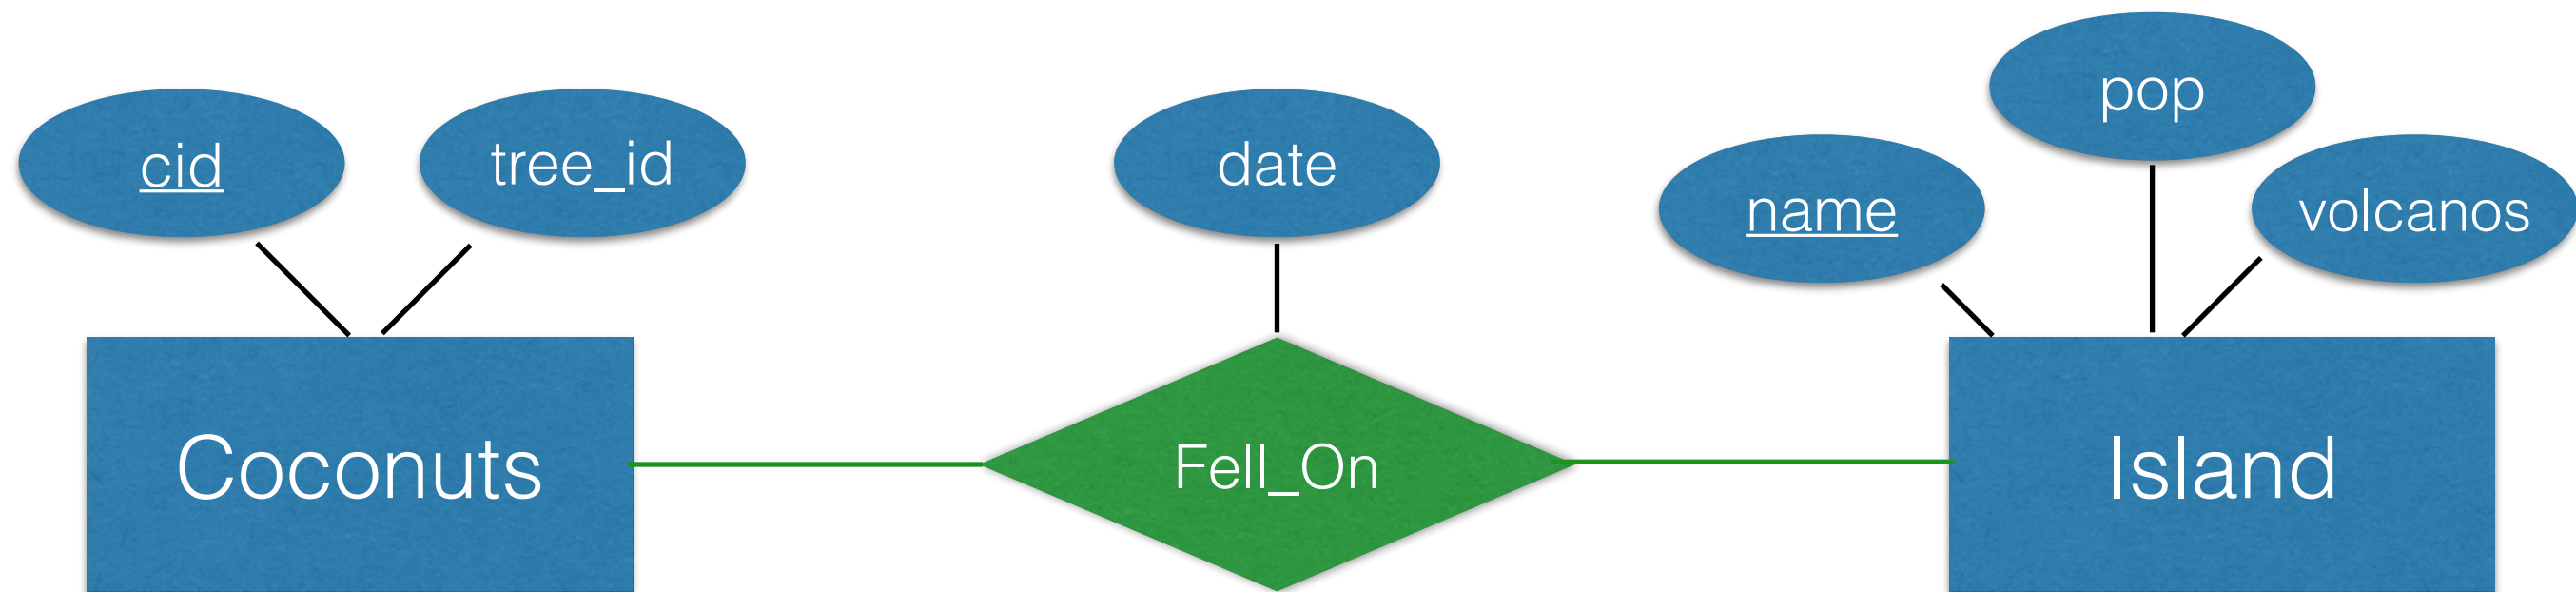
ER Diagrams: Quick Example



“Coconut participates in the Fell_On relationship 1 time.”
???

“A coconut falls on an island exactly once.”

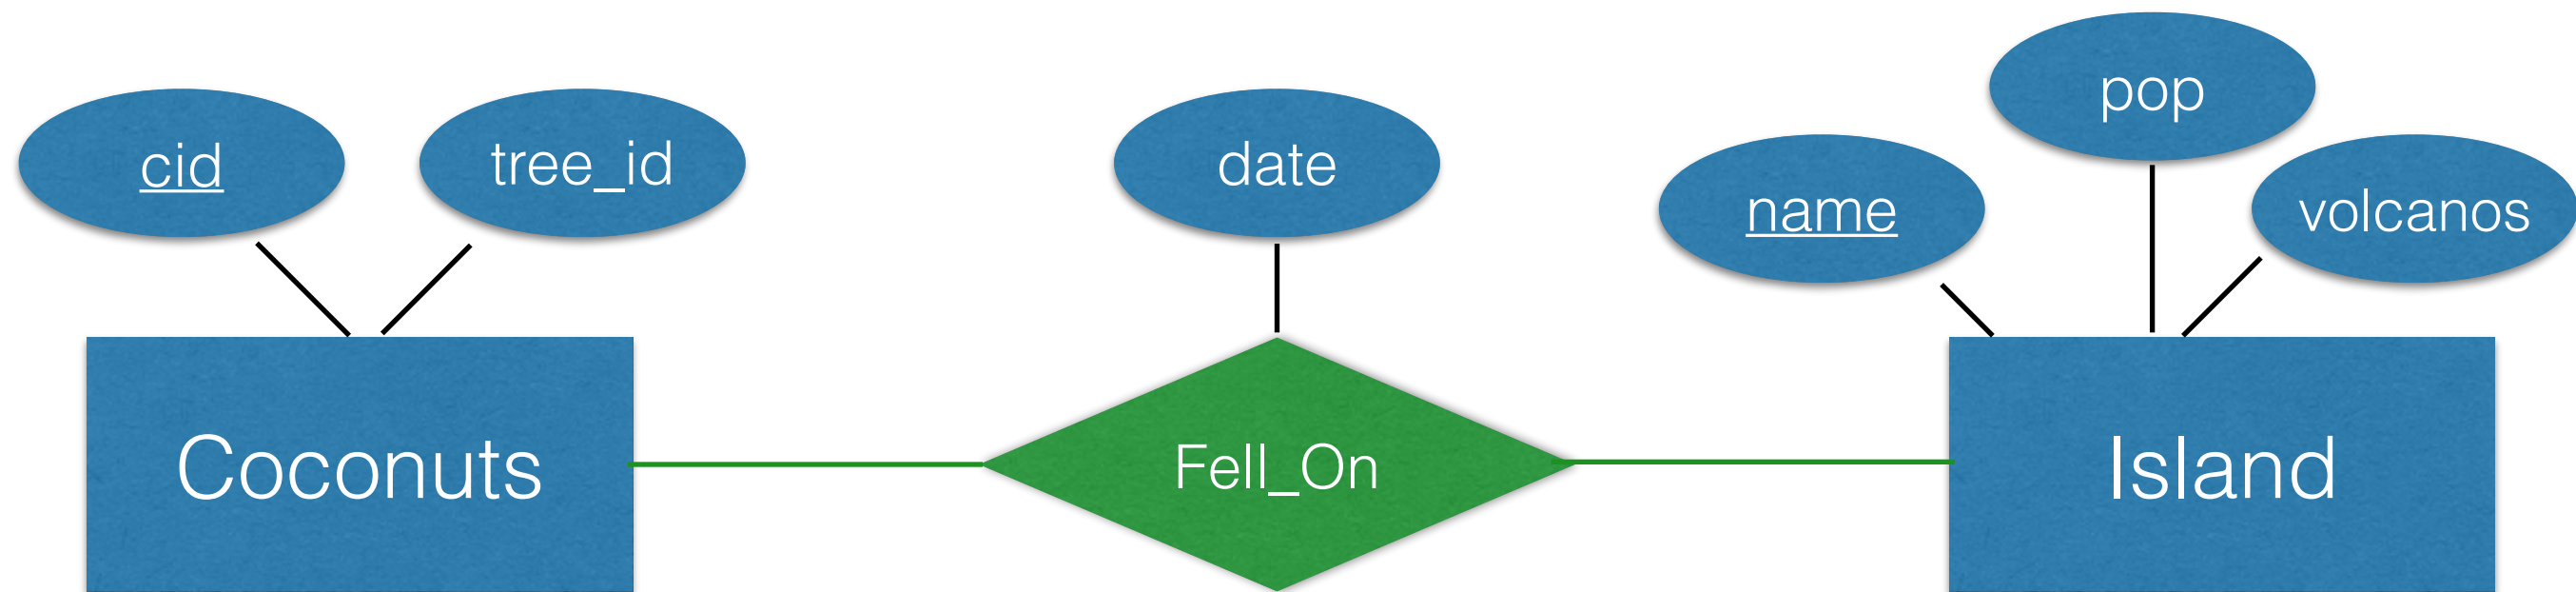
ER Diagrams: Quick Exercise



What constraint do we want from Coconuts?

- A. A coconut falls 0 or more times.
- B. A coconut falls 1 or more times
- C. A coconut falls 0 or 1 times.
- D. A coconut falls exactly once.

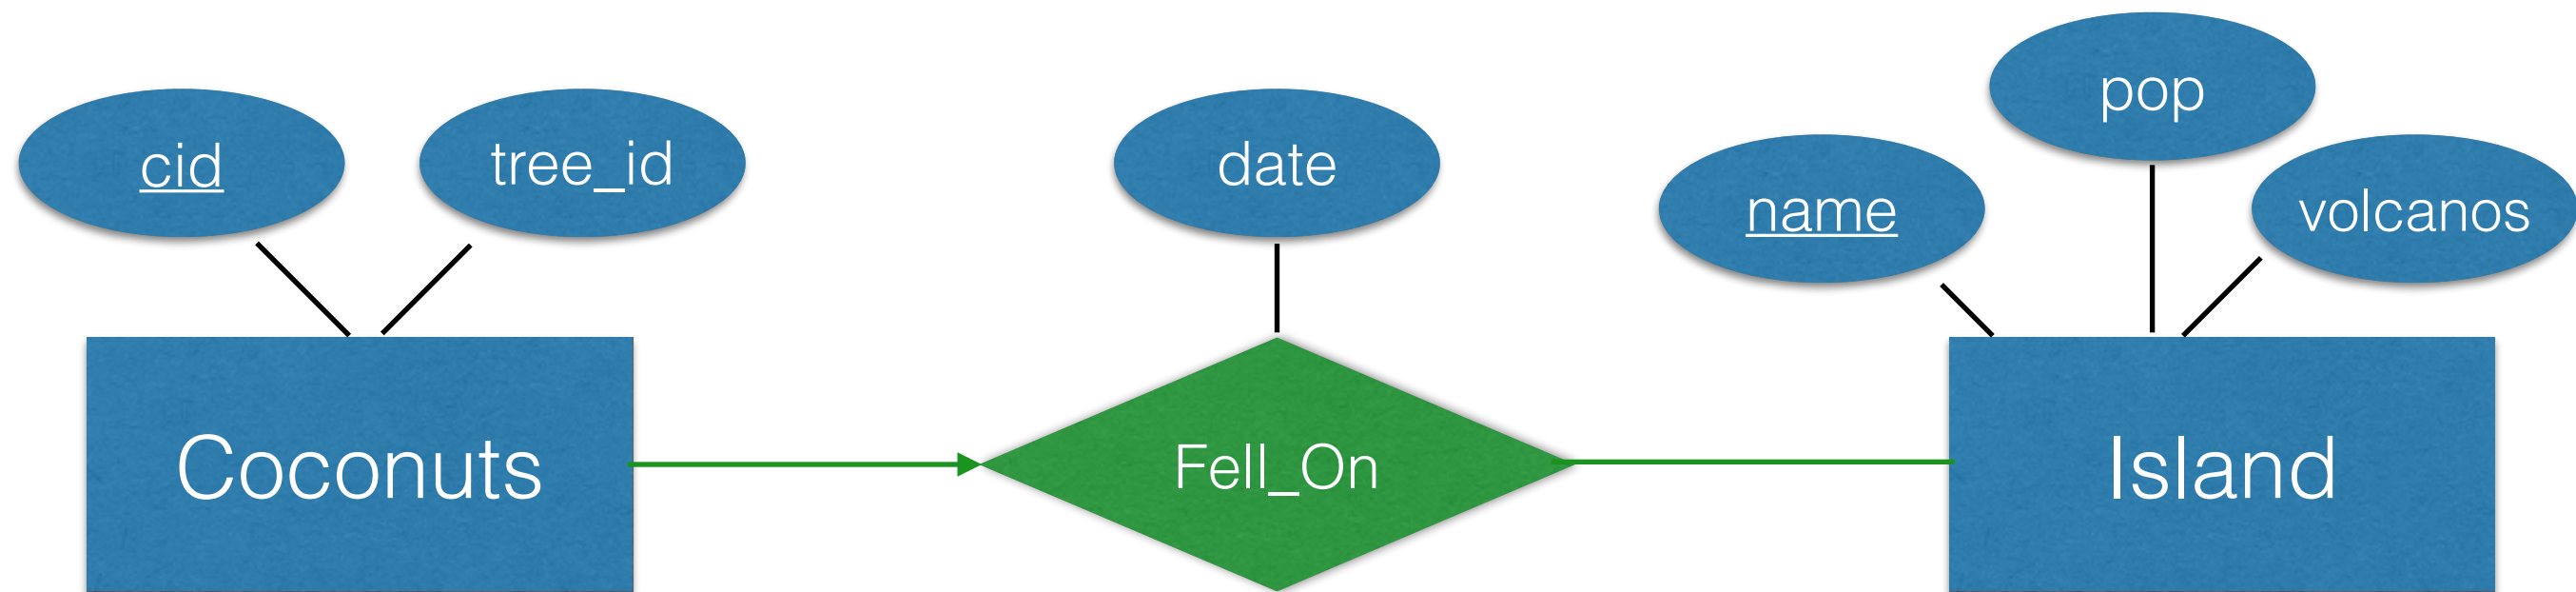
ER Diagrams: Quick Exercise



What constraint do we want from Coconuts?

- A. A coconut falls 0 or more times.
- B. A coconut falls 1 or more times
- C. A coconut falls 0 or 1 times.**
- D. A coconut falls exactly once.

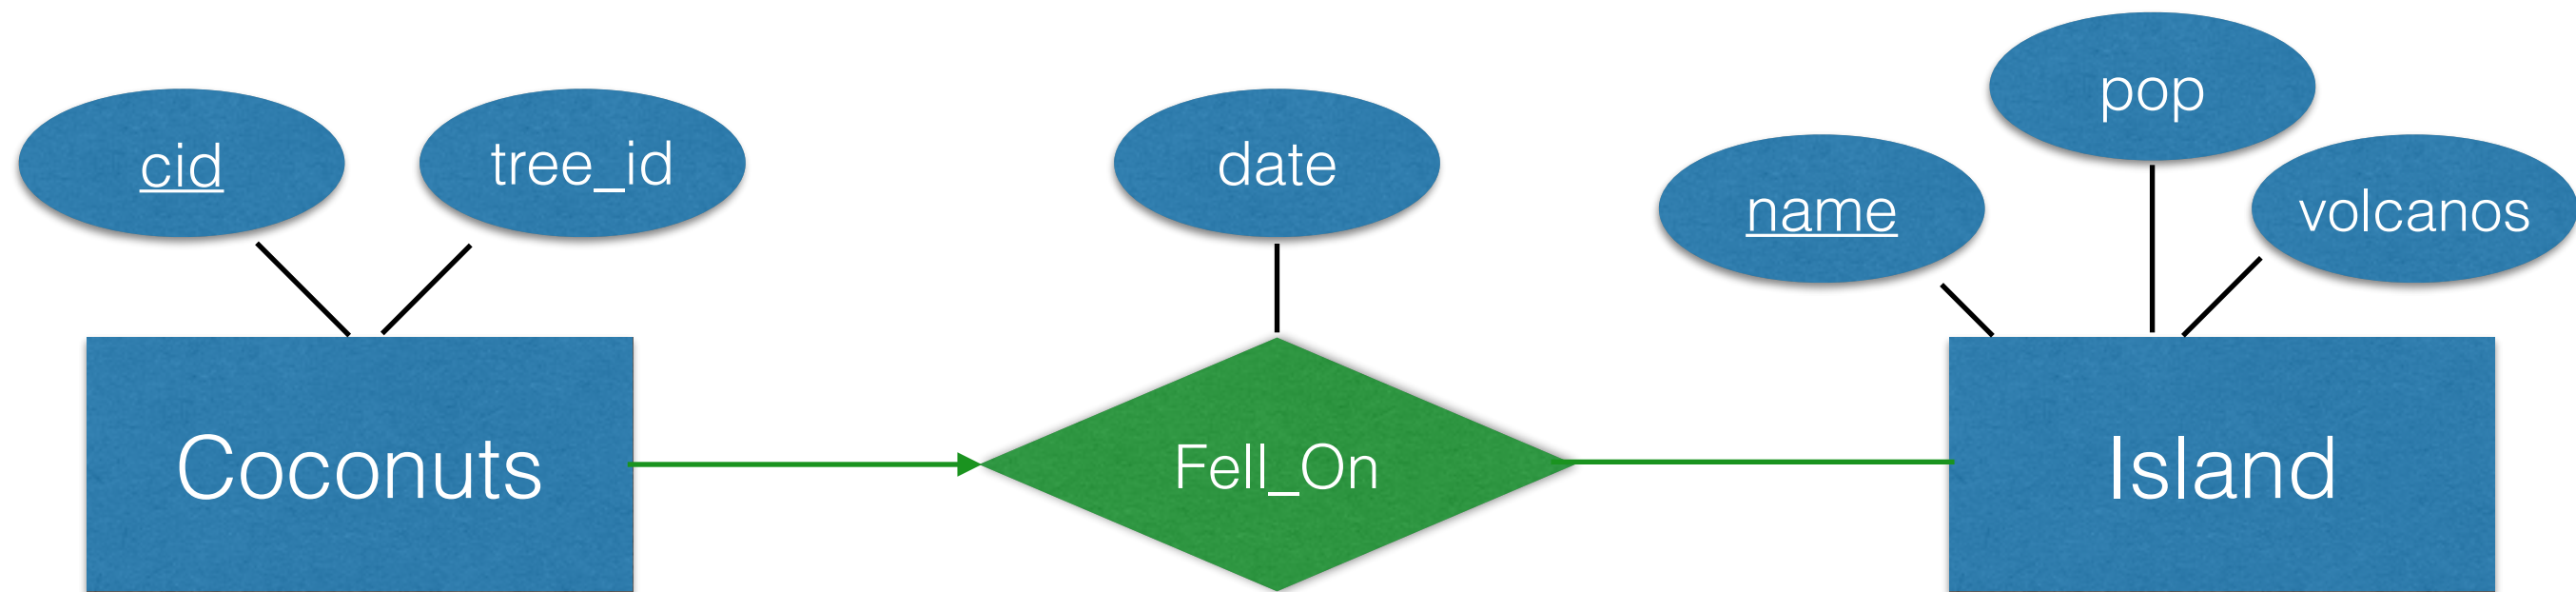
ER Diagrams: Quick Exercise



What constraint do we want from Island?

- A. A coconut falls 0 or more times.
- B. A coconut falls 1 or more times
- C. A coconut falls 0 or 1 times.
- D. A coconut falls exactly once.

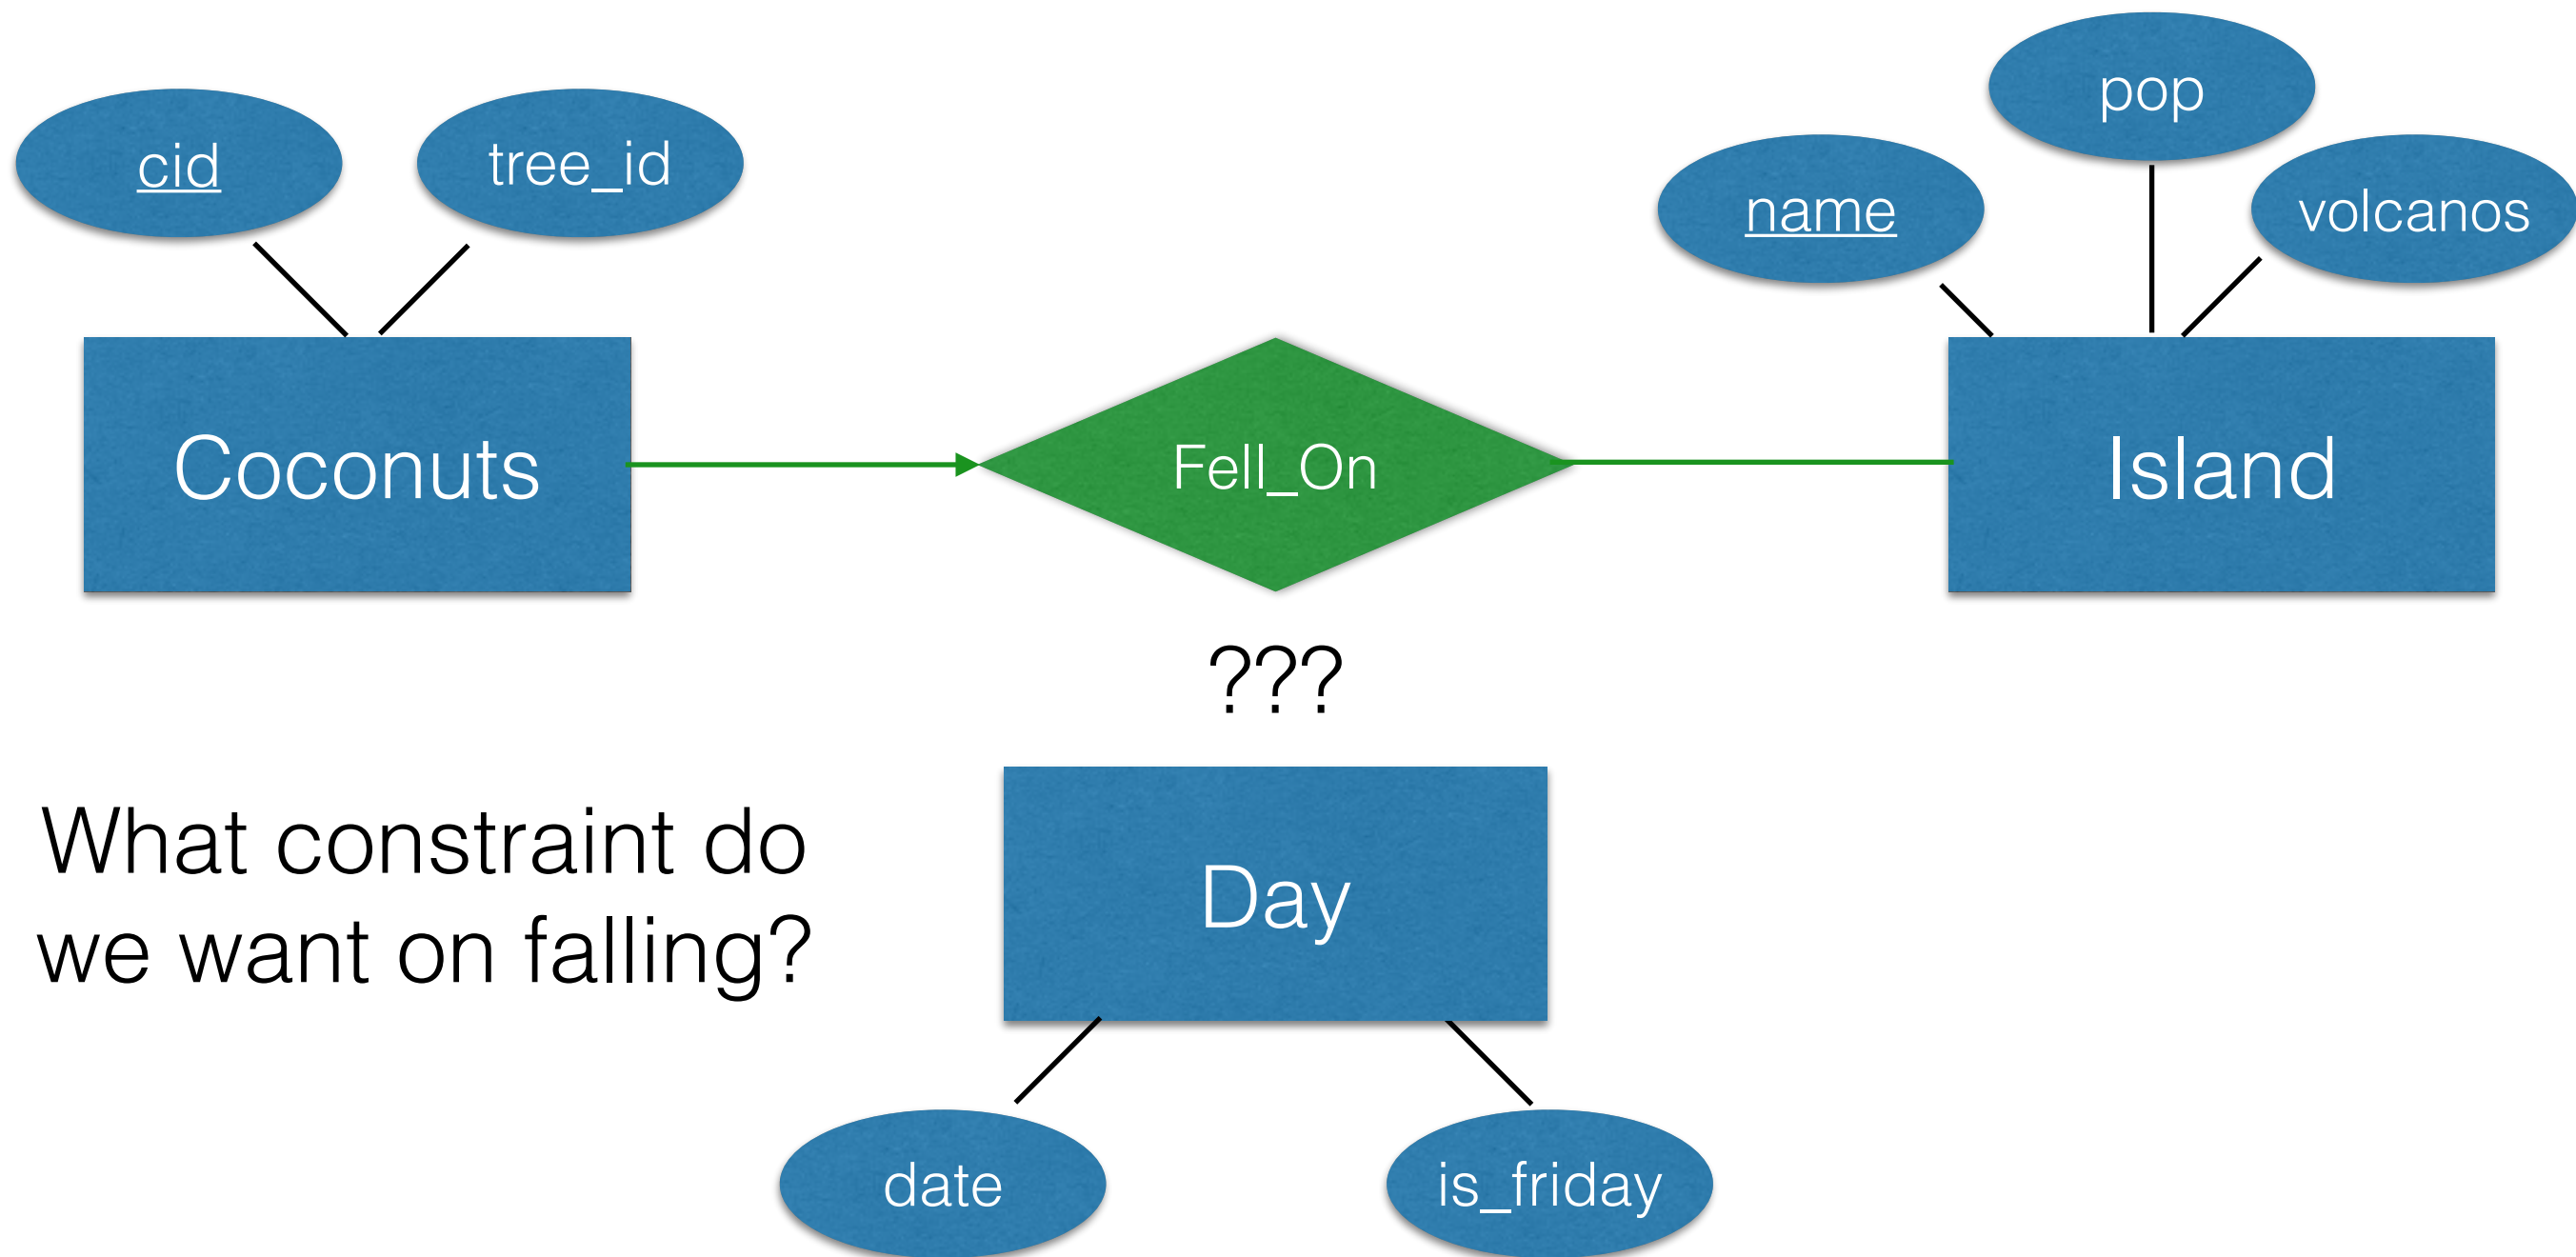
ER Diagrams: Quick Exercise



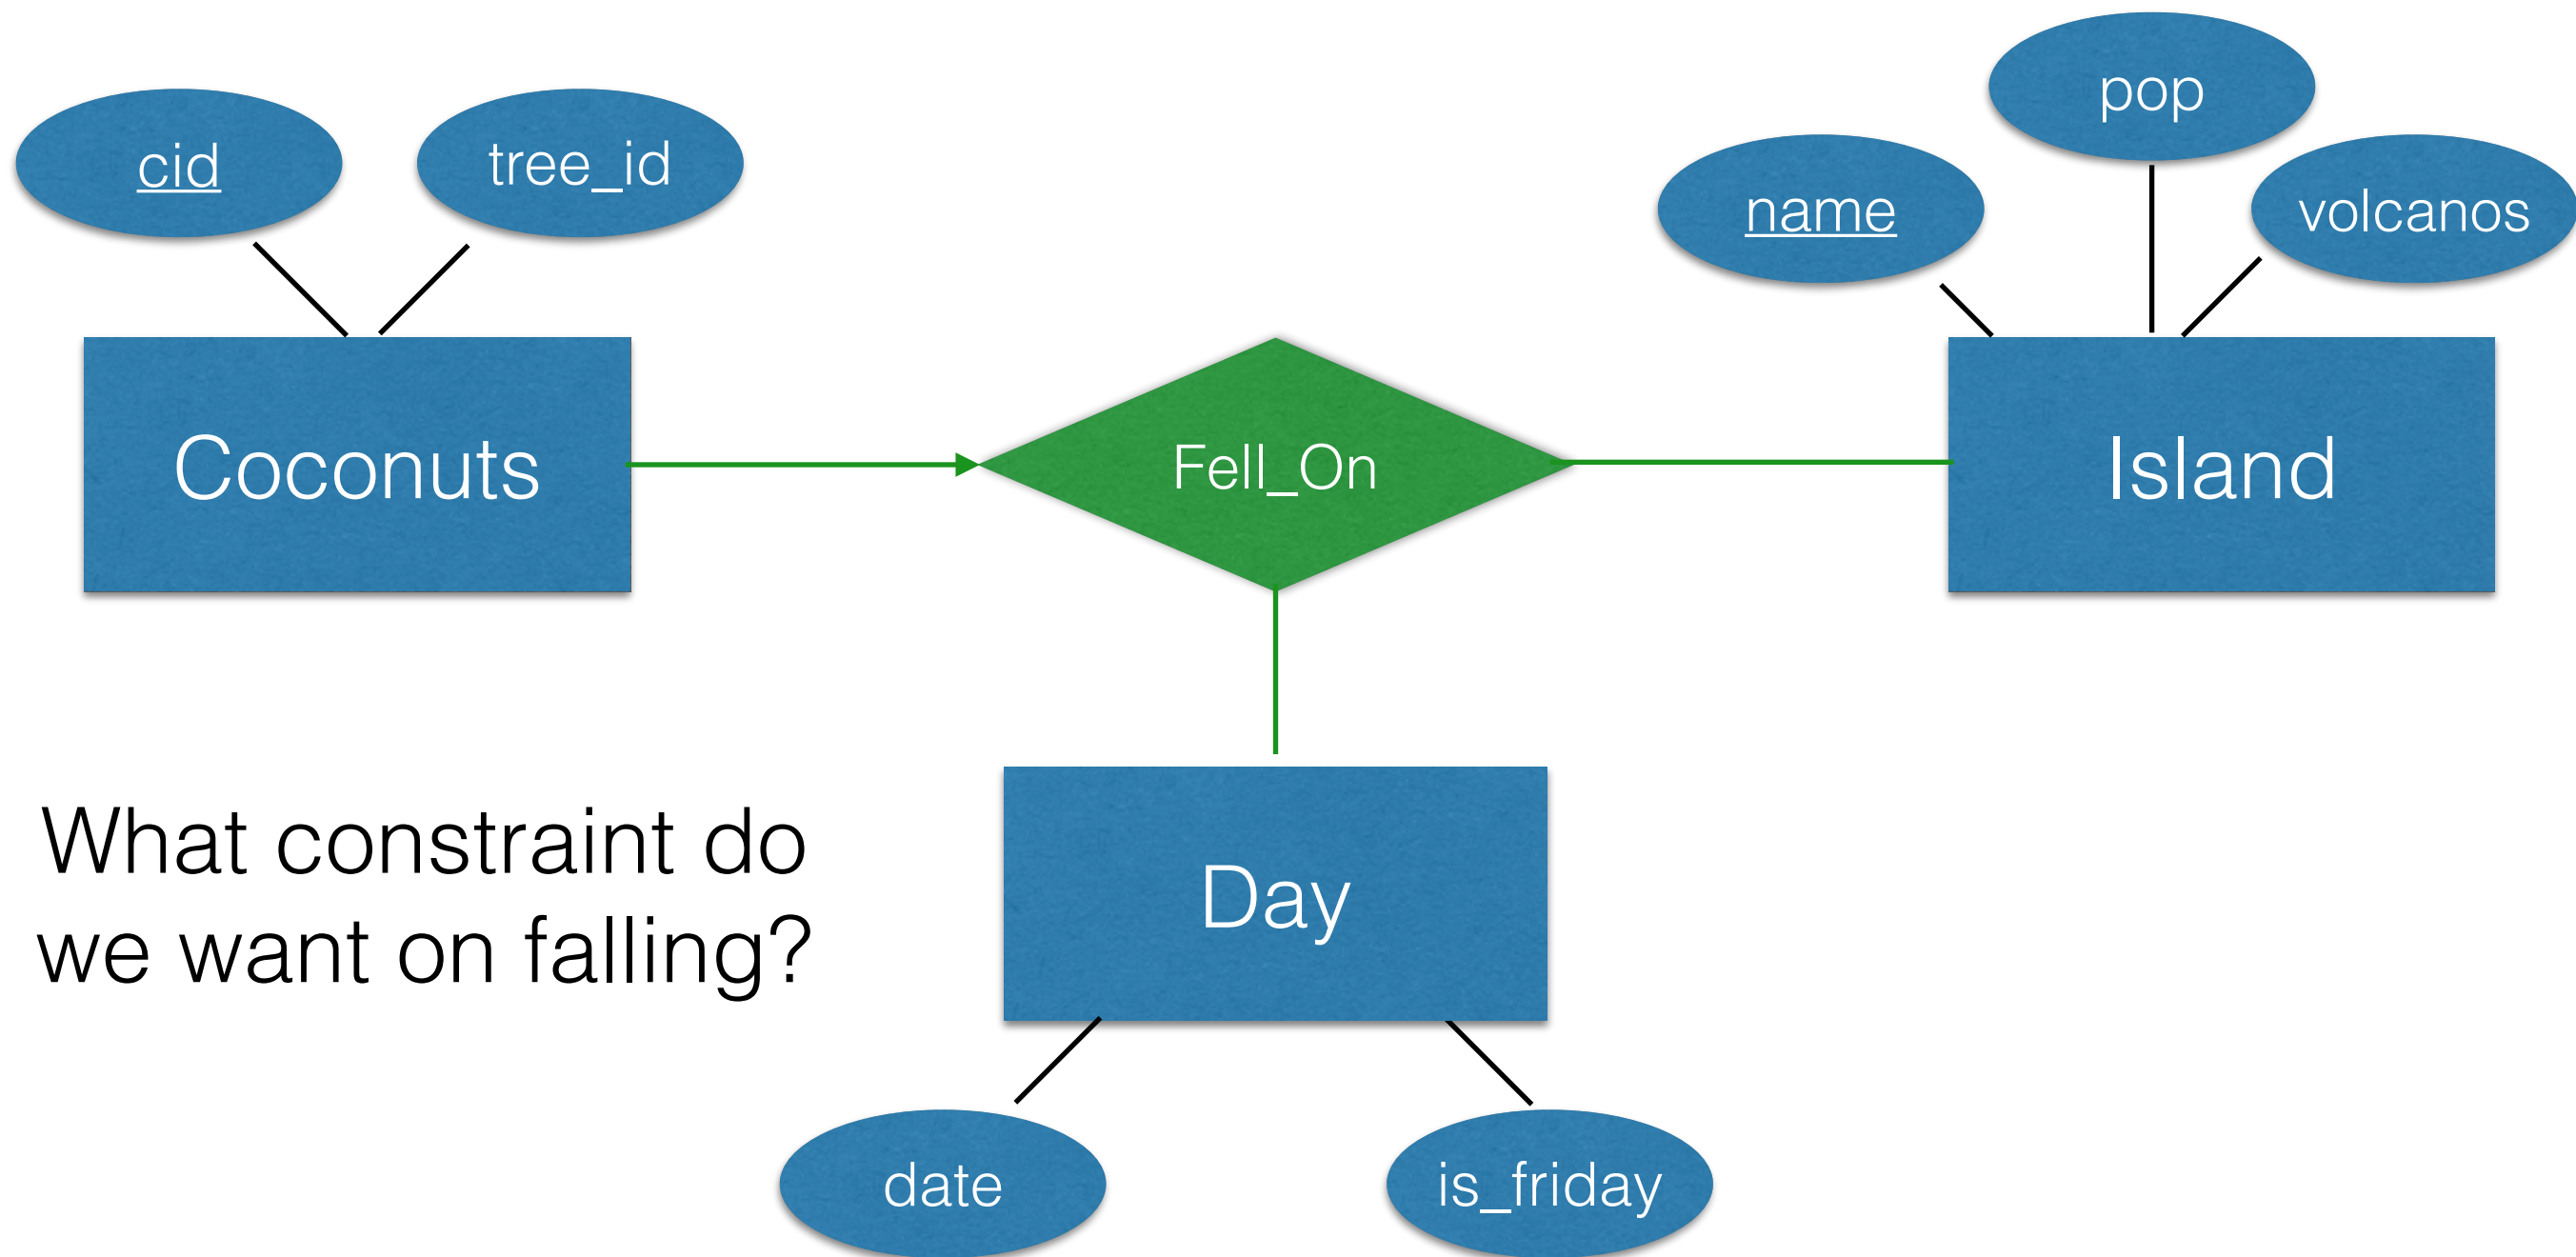
What constraint do we want from Island?

- A. A coconut falls 0 or more times.**
- B. A coconut falls 1 or more times
- C. A coconut falls 0 or 1 times.
- D. A coconut falls exactly once.

ER Diagrams: Ternary Relations

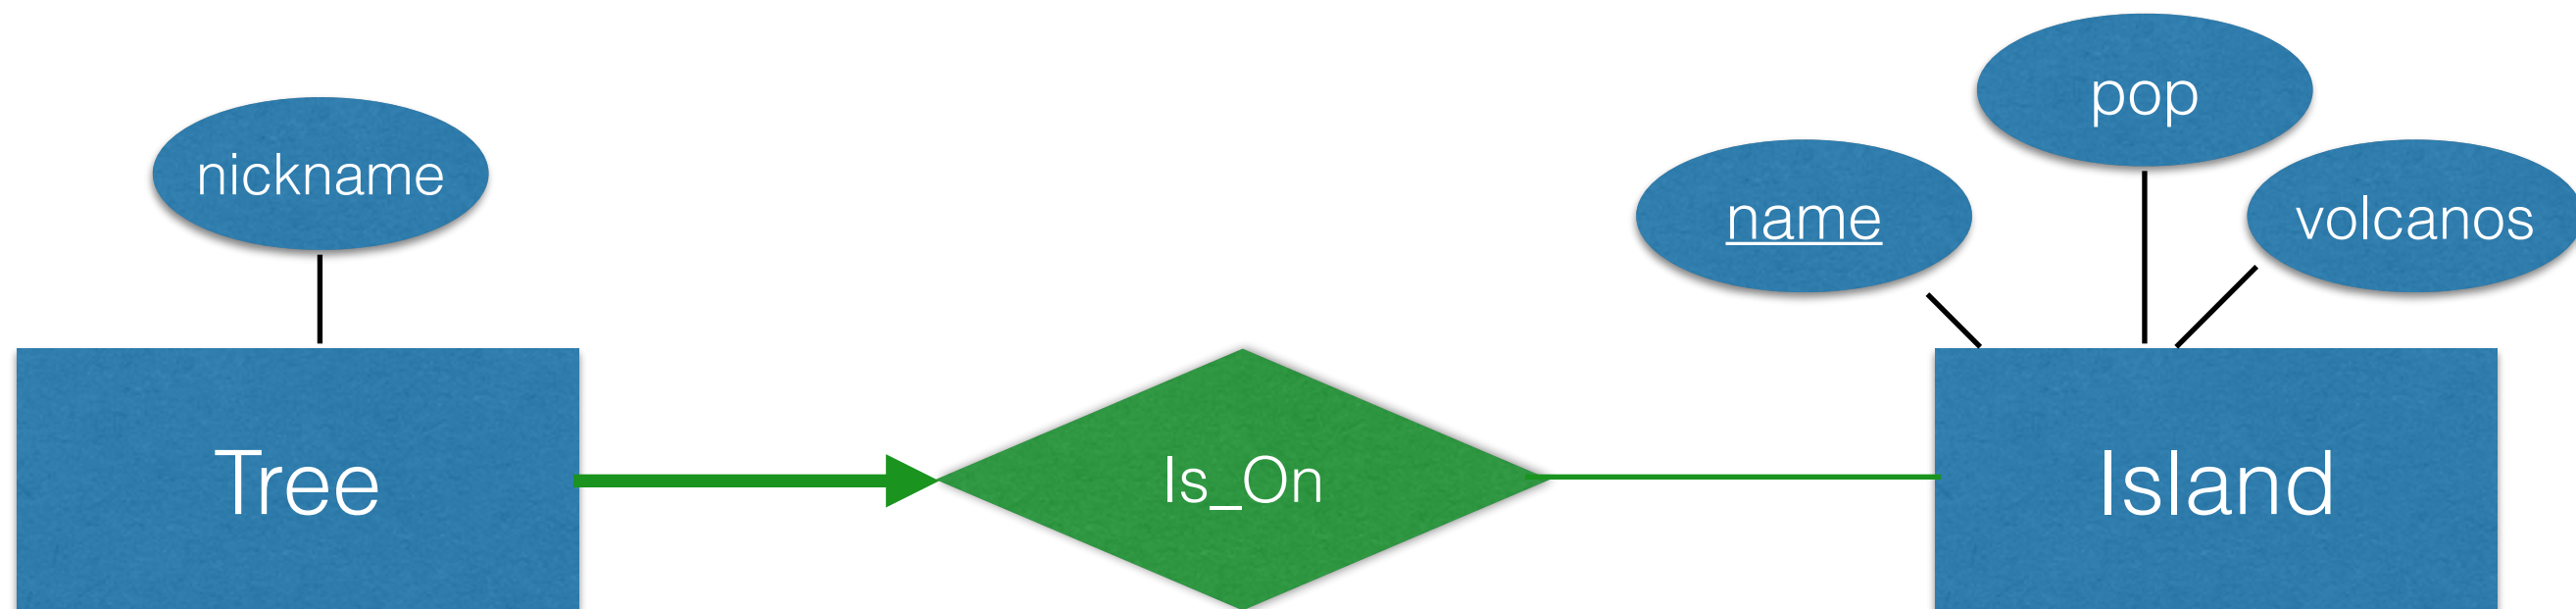


ER Diagrams: Ternary Relations



ER Diagrams: Weak Entities

- An entity that only makes sense in the context of a parent.



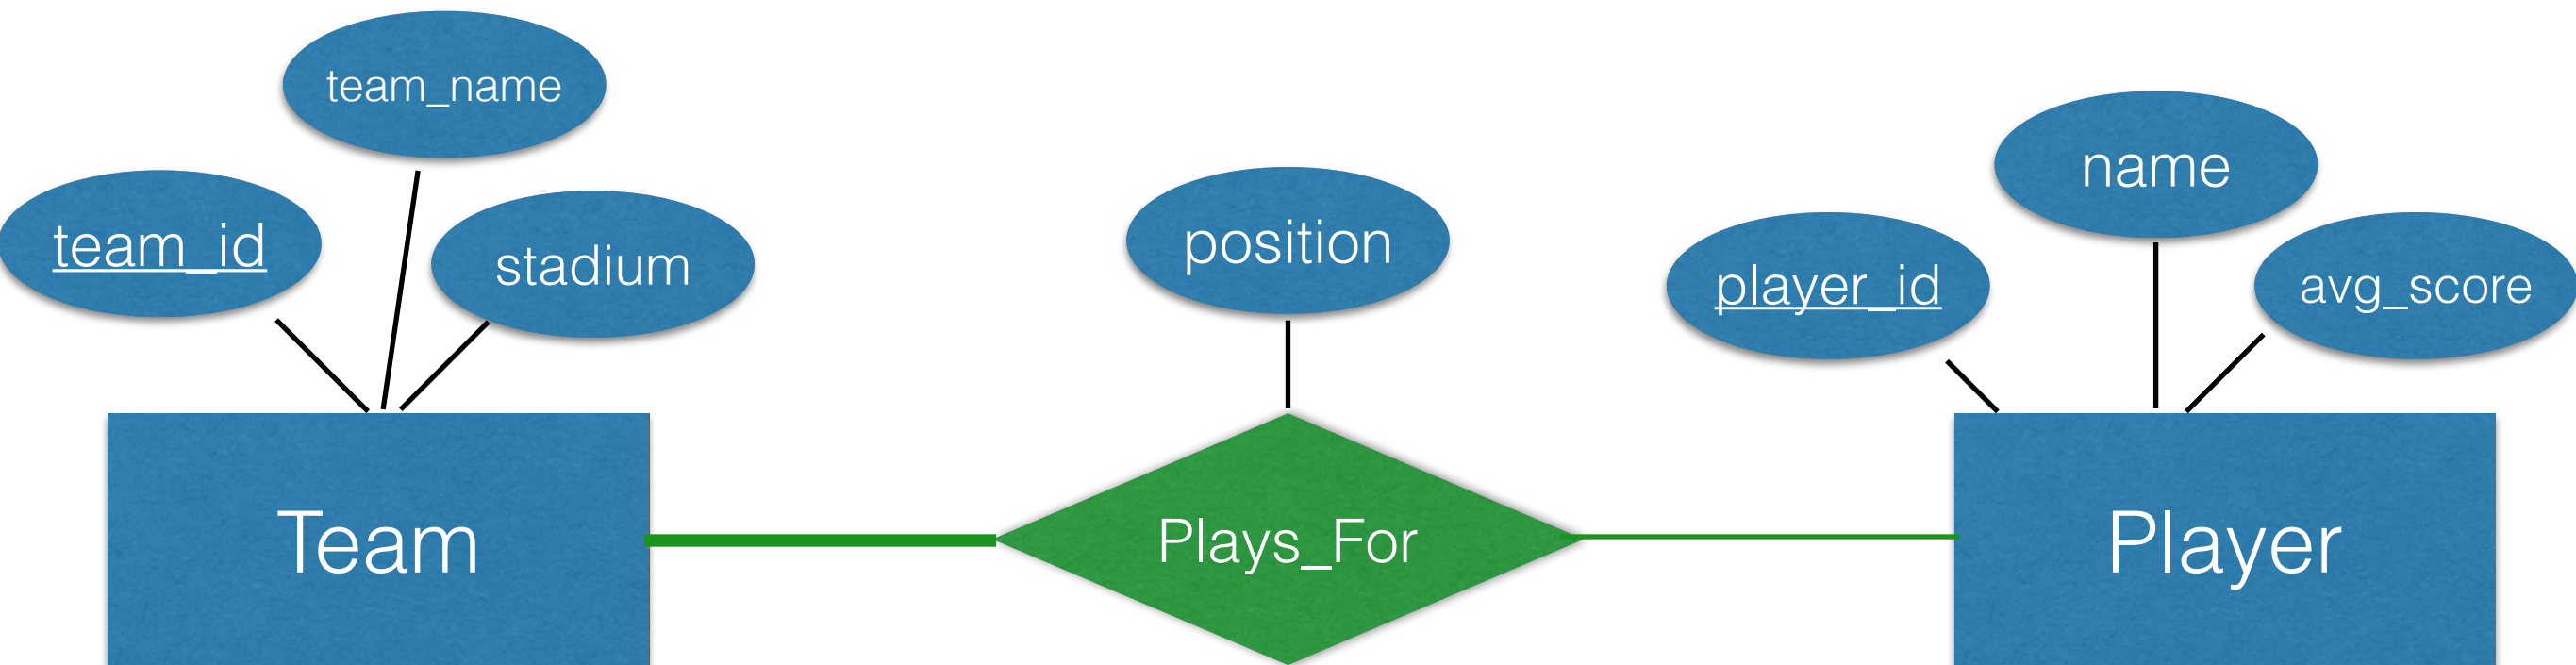
- A tree's key is (Island.name, Tree.nickname).
- There can be two trees with the same name, but not on the same island!
- A tree cannot exist without an island.

ER Diagrams: Exercise

We want to store sports teams and their players in our database. So let's first make an ER diagram! Every **Team** in our database will have a **team_id**, a **team_name**, and a **stadium** where they play their games. Each **Player** will have a **player_id**, **name** and their **avg_score**. Finally our database will contain who **Plays_For** which team and also the **position** that the player plays in.

Assume that a player can play for more than team, and every team must have at least one player. Draw an ER diagram for this database.

ER Diagram: Exercise

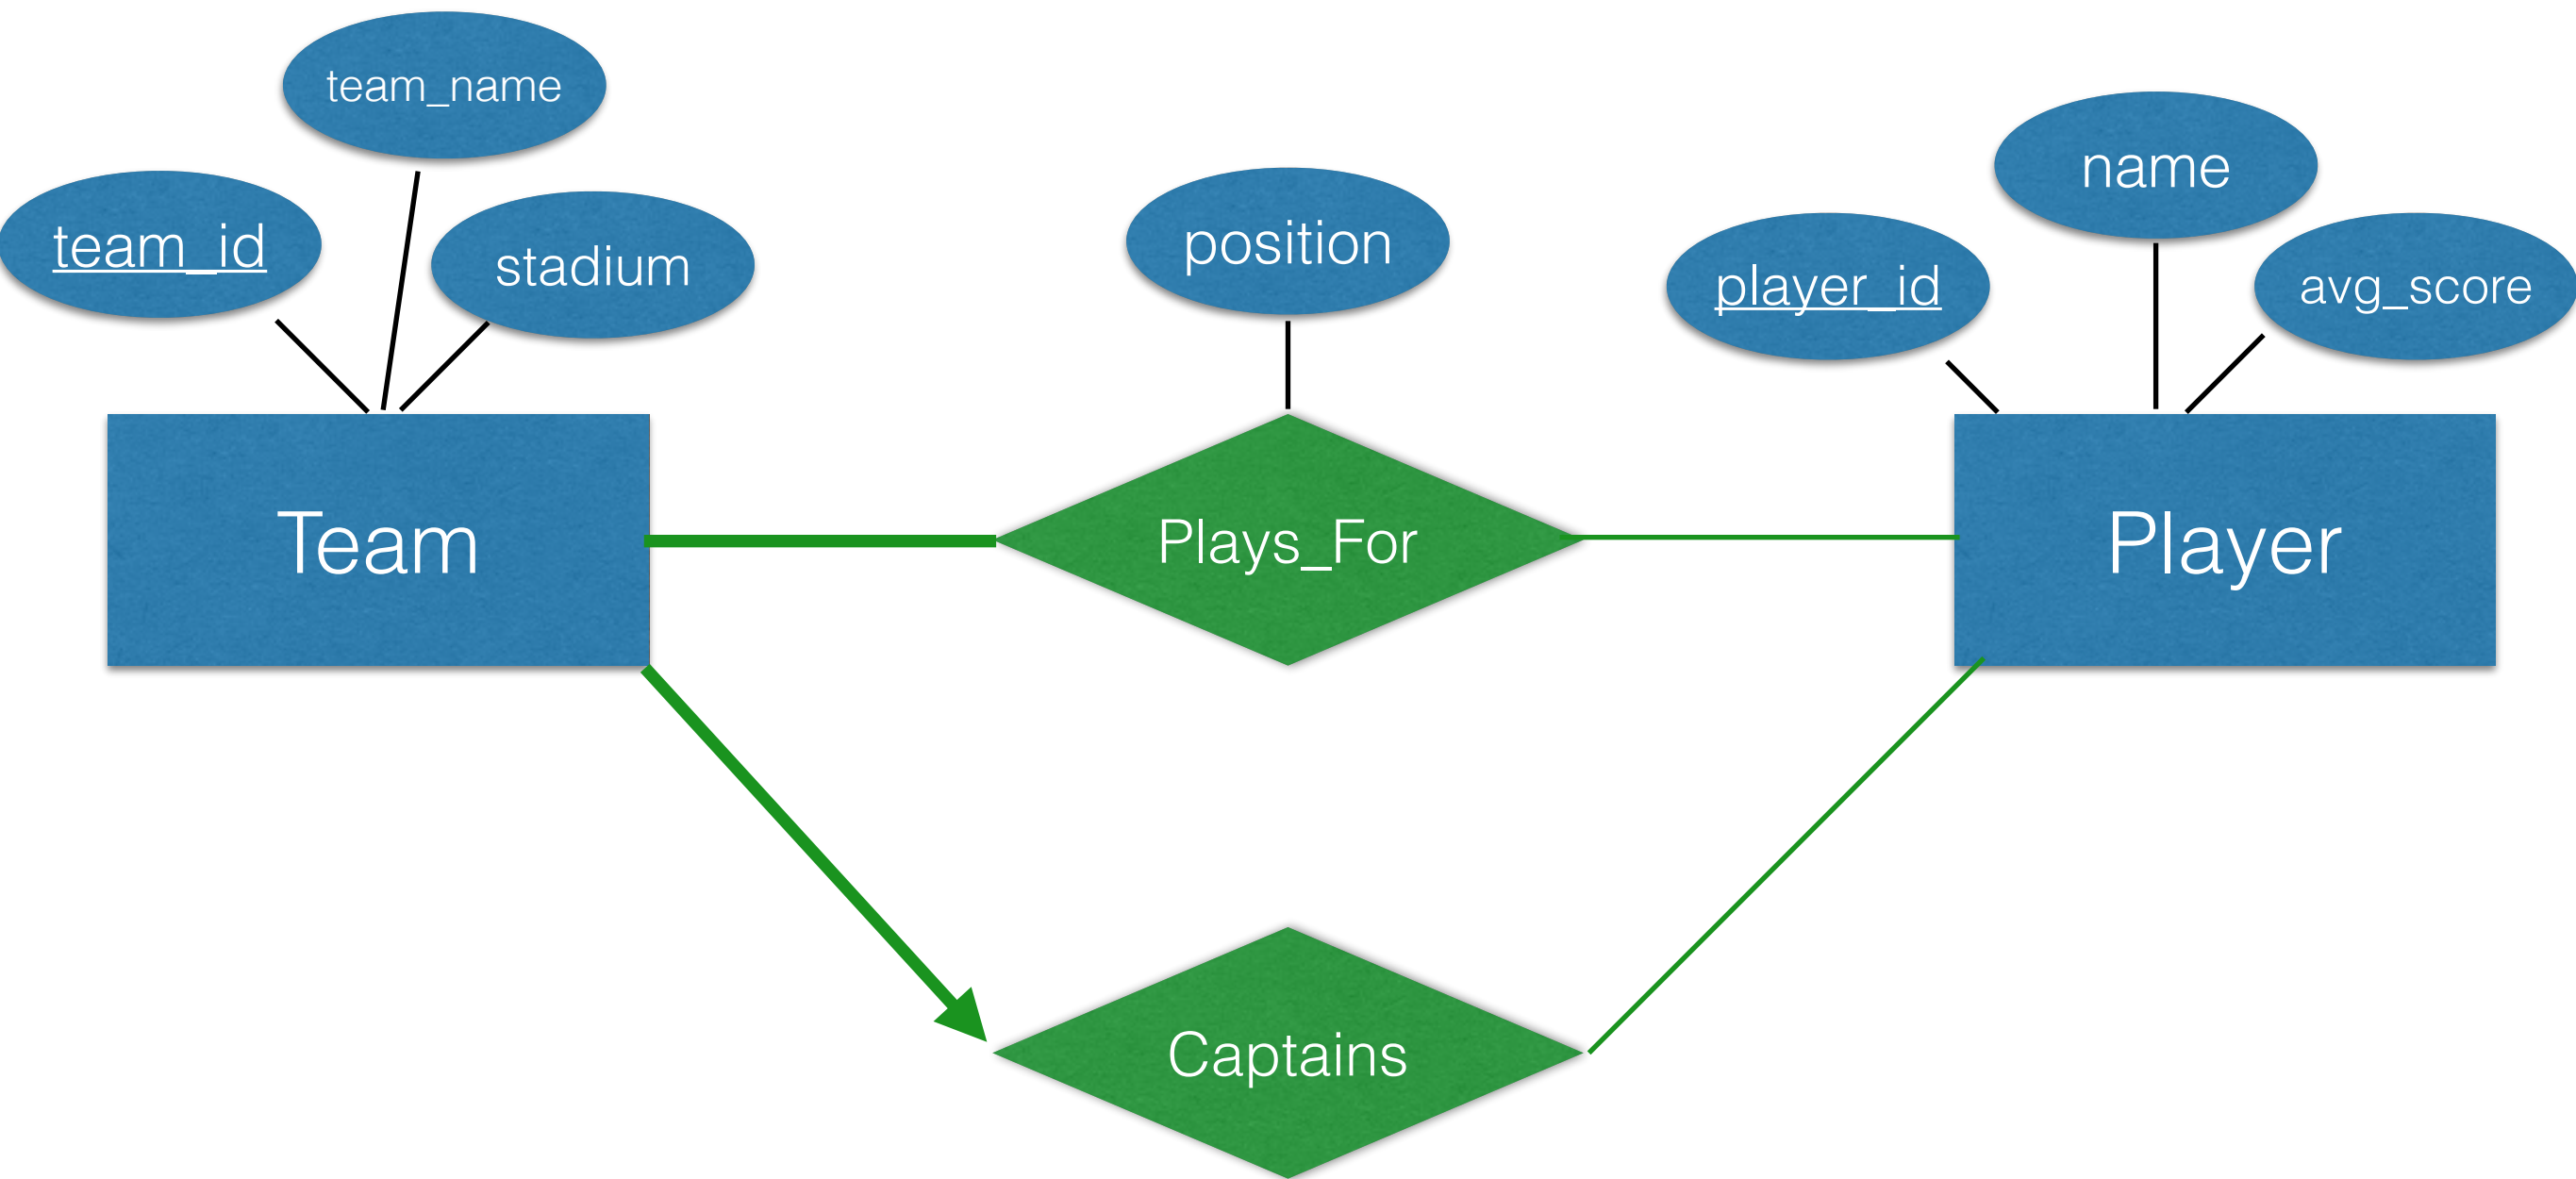


ER Diagrams: Exercise

We want to store sports teams and their players in our database. So let's first make an ER diagram! Every **Team** in our database will have a **team_id**, a **team_name**, and a **stadium** where they play their games. Each **Player** will have a **player_id**, **name** and their **avg_score**. Finally our database will contain who **Plays_For** which team and also the **position** that the player plays in.

Now let's say we also want to have a **Captain** for each time. A **Captain** is a **Player** and every team must have one **Captain**. (The constraints from the previous problem apply.) Redraw the ER diagram.

ER Diagram: Exercise



Advanced SQL

I'm sure you guys are more than sick of SQL, so if you don't want to hear this, feel free to leave.

Full SQL Query

```
SELECT [DISTINCT] <columns>  
FROM <tables>  
WHERE <predicate>  
[GROUP BY <grouping column>  
[HAVING <group filter>]  
[ORDER BY <order>]  
[LIMIT <limit>];
```

Combining Results of Queries

- **UNION** — combination of all the results of the first query and the second query.
- **INTERSECT** — tuples that are in both the first and the second query.
- **EXCEPT** — tuples that are in the first query but not in the second query.

SQL Example

Find `sids` of sailors who've reserved a red or a green boat.

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND  
      (B.color='red' or  
       B.color='green');
```


SQL Example: UNION

Find `sids` of sailors who've reserved a red or a green boat.

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND
      B.color='red'
UNION
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND
      B.color='green';
```

SQL Example

Find `sids` of sailors who've reserved a red and a green boat.

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND      ← WTF?  
      (B.color='red' and  
       B.color='green');
```

SQL Example: INTERSECT

Find `sids` of sailors who've reserved a red and a green boat.

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND
      B.color='red'
INTERSECT
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND
      B.color='green';
```

SQL Example: Self-Join

Find `sids` of sailors who've reserved a red and a green boat using a self-join!

```
SELECT R.sid
FROM Boats B1, Reserves R1,
      Boats B2, Reserves R2
WHERE R1.sid=R2.sid
      AND R1.bid=B1.bid
      AND R2.bid=B2.bid
      AND(B1.color='red'
          AND B2.color='green');
```

SQL Example: EXCEPT

Find `sids` of sailors who have not reserved at a boat.

```
SELECT S.sid  
FROM Sailors S  
EXCEPT  
SELECT S.sid  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid;
```

SQL Example: IN

Find `sids` of sailors who have reserved boat #102.

```
SELECT S.sid  
FROM Sailors S  
WHERE S.sid IN  
      (SELECT R.sid  
       FROM Reserves R  
       WHERE R.bid=102;
```

SQL Example: NOT IN


Find `sids` of sailors who have not reserved boat #102.

```
SELECT S.sid
FROM Sailors S
WHERE S.sid NOT IN
      (SELECT R.sid
       FROM Reserves R
       WHERE R.bid=102;
```

SQL Example: Correlated Subqueries

Find `sids` of sailors who have not reserved boat #102.

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS  
    (SELECT *  
     FROM Reserves R  
     WHERE R.bid=102 AND S.sid=R.sid;
```



SQL Joins

- **INNER:** Return only rows that match.
- **LEFT:** Return rows that match plus all unmatched rows from the table on the left side of the join clause.
- **RIGHT:** Return rows that match plus all unmatched rows from the table on the right side of the join clause.
- **OUTER:** Return rows that match plus all unmatched rows from the table on both sides of the join clause.

SQL: Left Outer Join

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/2014
95	103	11/12/2014

s.sid	s.sname	r.bid
22	Dustin	101
31	Lubber	103
95	Bob	

SQL: Right Outer Join

sid	bid	day
22	101	10/10/2014
95	103	11/12/2014

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

SQL: Full Outer Join

sid	bid	day
22	101	10/10/2014
95	103	11/12/2014

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Note that this is the same as the ROJ because **bid** is a foreign key in **Reserves**.

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

SQL Exercises

Songs (song_id, song_name, album_id, weeks_in_top_40)

Artists(artist_id, artist_name, first_year_active)

Albums (album_id, album_name, artist_id, year_released, genre)

Find all the album_ids and names for every artist who became active 2000 or later. If an artist does not have any albums, you should still include the artist in your output.

```
SELECT Ar.artist_id, Ar.artist_name,  
       Al.album_id, Al.album_name  
FROM Artists AR  
      LEFT OUTER JOIN Albums Al  
ON Ar.artist_id = Al.artist_id  
WHERE Ar.first_year_active >= 2000
```

SQL Exercises

Songs (song_id, song_name, album_id, weeks_in_top_40)

Artists(artist_id, artist_name, first_year_active)

Albums (album_id, album_name, artist_id, year_released, genre)

Find the song_ids and names for each song released that the artist of the album was active.

```
SELECT S.song_id, S.song_name
       Ar.artist_id, Ar.artist_name
FROM Song S INNER JOIN Albums Al
ON S.album_id=Al.album_id
INNER JOIN Song S
ON Al.artist_id=Ar.artist_id
AND Al.year_released=Ar.first_year_active
```